

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, roc_auc_score, confusion_matrix,
                             classification_report)

import xgboost as xgb
import lightgbm as lgb
```

```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("fedesoriano/heart-failure-prediction")

print("Path to dataset files:", path)
```

🔗 Path to dataset files: /kaggle/input/heart-failure-prediction

Double-click (or enter) to edit

```
df = pd.read_csv("/kaggle/input/heart-failure-prediction/heart.csv")
```

```
print(f"Dataset shape: {df.shape}")
```

🔗 Dataset shape: (918, 12)

```
print("\nFirst 5 rows:")
print(df.head())
```

🔗

First 5 rows:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

```
print("\nData types and missing values:")
print(df.info())
```

🔗

Data types and missing values:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Age                 918 non-null   int64
1   Sex                 918 non-null   object
2   ChestPainType       918 non-null   object
3   RestingBP           918 non-null   int64
4   Cholesterol         918 non-null   int64
5   FastingBS           918 non-null   int64
6   RestingECG          918 non-null   object
7   MaxHR              918 non-null   int64
8   ExerciseAngina      918 non-null   object
9   Oldpeak             918 non-null   float64
10  ST_Slope            918 non-null   object
11  HeartDisease        918 non-null   int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
None
```

```
print("\nSummary statistics:")
print(df.describe(include='all'))
```




Summary statistics:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	\
count	918.000000	918	918	918.000000	918.000000	918.000000	
unique	NaN	2	4	NaN	NaN	NaN	
top	NaN	M	ASY	NaN	NaN	NaN	
freq	NaN	725	496	NaN	NaN	NaN	
mean	53.510893	NaN	NaN	132.396514	198.799564	0.233115	
std	9.432617	NaN	NaN	18.514154	109.384145	0.423046	
min	28.000000	NaN	NaN	0.000000	0.000000	0.000000	
25%	47.000000	NaN	NaN	120.000000	173.250000	0.000000	
50%	54.000000	NaN	NaN	130.000000	223.000000	0.000000	
75%	60.000000	NaN	NaN	140.000000	267.000000	0.000000	
max	77.000000	NaN	NaN	200.000000	603.000000	1.000000	


	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	\
count	918	918.000000	918	918.000000	918	
unique	3	NaN	2	NaN	3	
top	Normal	NaN	N	NaN	Flat	
freq	552	NaN	547	NaN	460	
mean	NaN	136.809368	NaN	0.887364	NaN	
std	NaN	25.460334	NaN	1.066570	NaN	
min	NaN	60.000000	NaN	-2.600000	NaN	
25%	NaN	120.000000	NaN	0.000000	NaN	
50%	NaN	138.000000	NaN	0.600000	NaN	
75%	NaN	156.000000	NaN	1.500000	NaN	
max	NaN	202.000000	NaN	6.200000	NaN	

	HeartDisease
count	918.000000
unique	NaN
top	NaN
freq	NaN
mean	0.553377
std	0.497414
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000


```
# 1. Target variable distribution
plt.figure(figsize=(6, 4)) # Increased figure size
plt.subplot(1, 1, 1) # Adjusted subplot for single plot
sns.countplot(x='HeartDisease', data=df)
plt.title('Heart Disease Distribution')
plt.show() # Added plt.show() to display plot immediately
```



```
# 1. Target variable distribution
plt.figure(figsize=(6, 4)) # Increased figure size
plt.subplot(1, 1, 1) # Adjusted subplot for single plot
sns.countplot(x='HeartDisease', data=df)
plt.title('Heart Disease Distribution')
plt.show() # Added plt.show() to display plot immediately
```




```
# 2. Age distribution
plt.figure(figsize=(6, 4)) # Increased figure size
plt.subplot(1, 1, 1) # Adjusted subplot for single plot
sns.histplot(df['Age'], bins=20, kde=True)
plt.title('Age Distribution')
plt.show() # Added plt.show() to display plot immediately
```



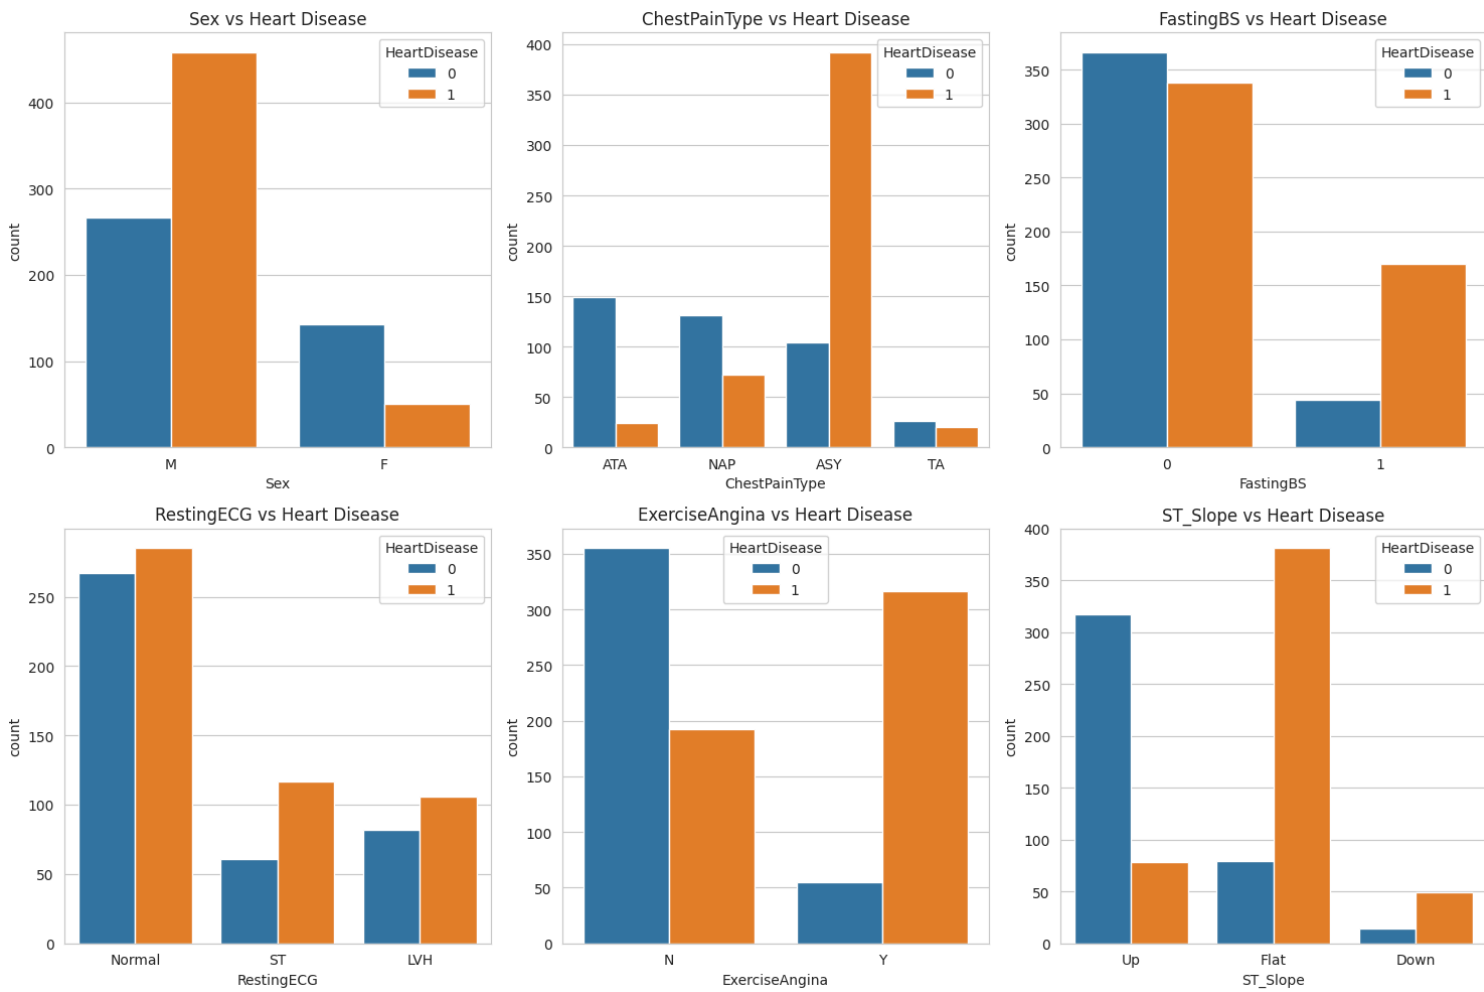
```
# 4. RestingBP distribution
plt.figure(figsize=(6, 4)) # Increased figure size
plt.subplot(1, 1, 1) # Adjusted subplot for single plot
sns.histplot(df['RestingBP'], bins=20, kde=True)
plt.title('Resting Blood Pressure Distribution')

plt.tight_layout()
plt.show()
```




```
# Categorical features visualization
cat_features = ['Sex', 'ChestPainType', 'FastingBS', 'RestingECG', 'ExerciseAngina', 'ST_Slope']
```

```
plt.figure(figsize=(15, 10))
for i, feature in enumerate(cat_features, 1):
    plt.subplot(2, 3, i)
    sns.countplot(x=feature, hue='HeartDisease', data=df)
    plt.title(f'{feature} vs Heart Disease')
plt.tight_layout()
plt.show()
```



```
# Correlation matrix
plt.figure(figsize=(12, 8))
numeric_df = df.select_dtypes(include=['int64', 'float64'])
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Matrix')
plt.show()
```




```
# Check for missing values
print("Missing values per column:")
print(df.isnull().sum())
```

```
Missing values per column:
Age          0
Sex          0
ChestPainType  0
RestingBP    0
Cholesterol  0
FastingBS    0
RestingECG    0
MaxHR        0
ExerciseAngina  0
Oldpeak      0
ST_Slope     0
HeartDisease  0
dtype: int64
```

```
# Check for duplicates
print(f"\nNumber of duplicates: {df.duplicated().sum()}")
```

```
Number of duplicates: 0
```

```
# Check for outliers
numeric_features = ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak']
plt.figure(figsize=(15, 5))
for i, feature in enumerate(numeric_features, 1):
    plt.subplot(1, 5, i)
    sns.boxplot(y=df[feature])
    plt.title(feature)
plt.tight_layout()
plt.show()
```



```
# Investigate zero values in key features
print("\nPatients with zero RestingBP:", len(df[df['RestingBP'] == 0]))
print("Patients with zero Cholesterol:", len(df[df['Cholesterol'] == 0]))
```

```
Patients with zero RestingBP: 1
Patients with zero Cholesterol: 172
```

```
# Handle outliers and invalid values
# Replace zero values in RestingBP with median (zeros are likely errors)
df['RestingBP'] = df['RestingBP'].replace(0, df['RestingBP'].median())
```

```
# Cholesterol has many zeros - these might be valid (fasting) but we'll treat as missing
# Replace zero cholesterol with median of non-zero values
df['Cholesterol'] = df['Cholesterol'].replace(0, df[df['Cholesterol'] > 0]['Cholesterol'].median())
```

```
# Convert categorical features
df['Sex'] = df['Sex'].map({'M': 1, 'F': 0})
df['ExerciseAngina'] = df['ExerciseAngina'].map({'Y': 1, 'N': 0})
```

```
# One-hot encoding for other categorical features
df = pd.get_dummies(df, columns=['ChestPainType', 'RestingECG', 'ST_Slope'], drop_first=True)
```

```
# Verify the cleaned data
print("\nAfter cleaning:")
print(df.head())
print(df.info())
```

```
After cleaning:
   Age  Sex  RestingBP  Cholesterol  FastingBS  MaxHR  ExerciseAngina  \
0   40    1        140          289          0     172             0
1   49    0        160          180          0     156             0
2   37    1        130          283          0      98             0
3   48    0        138          214          0     108             1
4   54    1        150          195          0     122             0

   Oldpeak  HeartDisease  ChestPainType_AT  ChestPainType_NAP  \
0      0.0             0             True             False
1      1.0             1             False             True
2      0.0             0             True             False
3      1.5             1             False             False
4      0.0             0             False             True

   ChestPainType_TA  RestingECG_Normal  RestingECG_ST  ST_Slope_Flat  \
0             False             True             False             False
1             False             True             False             True
2             False             False             True             False
3             False             True             False             True
4             False             True             False             False

   ST_Slope_Up
0             True
1             False
2             True
3             False
4             True
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 16 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Age                 918 non-null   int64
 1   Sex                 918 non-null   int64
 2   RestingBP           918 non-null   int64
 3   Cholesterol          918 non-null   int64
 4   FastingBS           918 non-null   int64
 5   MaxHR               918 non-null   int64
 6   ExerciseAngina       918 non-null   int64
 7   Oldpeak             918 non-null   float64
 8   HeartDisease         918 non-null   int64
 9   ChestPainType_ATA    918 non-null   bool
10   ChestPainType_NAP    918 non-null   bool
11   ChestPainType_TA     918 non-null   bool
12   RestingECG_Normal    918 non-null   bool
13   RestingECG_ST        918 non-null   bool
14   ST_Slope_Flat        918 non-null   bool
15   ST_Slope_Up          918 non-null   bool
dtypes: bool(7), float64(1), int64(8)
memory usage: 71.0 KB
None
```

```
# Create age bins
df['AgeGroup'] = pd.cut(df['Age'], bins=[0, 40, 50, 60, 70, 100],
                        labels=['<40', '40-50', '50-60', '60-70', '70+'])
```

```
# Create HR/BP ratio
df['HR_BP_Ratio'] = df['MaxHR'] / df['RestingBP']
# Replace infinite values with a large number or the maximum finite value
df.replace([np.inf, -np.inf], np.nan, inplace=True)
df['HR_BP_Ratio'].fillna(df['HR_BP_Ratio'].max(), inplace=True)
```

```
# Create cholesterol to age ratio
df['Chol_Age_Ratio'] = df['Cholesterol'] / df['Age']
# Replace infinite values with a large number or the maximum finite value
df.replace([np.inf, -np.inf], np.nan, inplace=True)
df['Chol_Age_Ratio'].fillna(df['Chol_Age_Ratio'].max(), inplace=True)
```

```
# Verify new features
print("\nFeature engineering:")
print(df[['AgeGroup', 'HR_BP_Ratio', 'Chol_Age_Ratio']].head())
```



```
Feature engineering:
  AgeGroup  HR_BP_Ratio  Chol_Age_Ratio
0      <40      1.228571      7.225000
1     40-50      0.975000      3.673469
2      <40      0.753846      7.648649
3     40-50      0.782609      4.458333
4     50-60      0.813333      3.611111
```

```
# Split features and target
X = df.drop('HeartDisease', axis=1)
y = df['HeartDisease']
```

```
# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42,
                                                    stratify=y)
```

```
# Handle class imbalance with SMOTE
# Drop AgeGroup column as it's not numeric and causes error with SMOTE
X_train_numeric = X_train.drop('AgeGroup', axis=1)

smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train_numeric, y_train)

print("\nClass distribution before SMOTE:")
print(y_train.value_counts())
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_res).value_counts())
```



```
Class distribution before SMOTE:
HeartDisease
1      406
```

```
0 328
Name: count, dtype: int64
```

```
Class distribution after SMOTE:
HeartDisease
1    406
0    406
Name: count, dtype: int64
```

```
# Define numeric and categorical features for preprocessing
numeric_features = ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak',
                    'HR_BP_Ratio', 'Chol_Age_Ratio']
categorical_features = ['Sex', 'FastingBS', 'ExerciseAngina',
                        'ChestPainType_ATA', 'ChestPainType_NAP',
                        'ChestPainType_TA', 'RestingECG_Normal',
                        'RestingECG_ST', 'ST_Slope_Flat', 'ST_Slope_Up']
```

```
# Create preprocessing pipeline
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', 'passthrough', categorical_features)
    ])

```

```
# Apply preprocessing
X_train_preprocessed = preprocessor.fit_transform(X_train_res)
X_test_preprocessed = preprocessor.transform(X_test)
```

```
# XGBoost model
xgb_model = xgb.XGBClassifier(
    objective='binary:logistic',
    n_estimators=1000,
    learning_rate=0.01,
    max_depth=5,
    subsample=0.8,
    colsample_bytree=0.8,
    gamma=1,
    reg_alpha=0.1,
    reg_lambda=1,
    random_state=42,
    early_stopping_rounds=50,
    eval_metric='auc'
)

```

```
# Train the model
xgb_model.fit(
    X_train_preprocessed,
    y_train_res,
    eval_set=[(X_test_preprocessed, y_test)],
    verbose=10
)

```

```
0] validation_0-auc:0.87673
10] validation_0-auc:0.91129
20] validation_0-auc:0.90638
30] validation_0-auc:0.90644
40] validation_0-auc:0.90513
50] validation_0-auc:0.90561
52] validation_0-auc:0.90609
```

```
XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=0.8, device=None, early_stopping_rounds=50,
               enable_categorical=False, eval_metric='auc', feature_types=None,
               feature_weights=None, gamma=1, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=0.01, max_bin=None, max_cat_threshold=None,
               max_cat_to_onehot=None, max_delta_step=None, max_depth=5,
               max_leaves=None, min_child_weight=None, missing=nan,
               monotone_constraints=None, multi_strategy=None, n_estimators=1000,
               n_jobs=None, num_parallel_tree=None, ...)
```

```
# Make predictions
xgb_preds = xgb_model.predict(X_test_preprocessed)
xgb_probs = xgb_model.predict_proba(X_test_preprocessed)[:, 1]
```




```
# Evaluate performance
print("\nXGBoost Performance:")
print(f"Accuracy: {accuracy_score(y_test, xgb_preds):.4f}")
print(f"Precision: {precision_score(y_test, xgb_preds):.4f}")
print(f"Recall: {recall_score(y_test, xgb_preds):.4f}")
print(f"F1 Score: {f1_score(y_test, xgb_preds):.4f}")
print(f"ROC AUC: {roc_auc_score(y_test, xgb_probs):.4f}")
```




```
XGBoost Performance:
Accuracy: 0.8587
Precision: 0.8958
Recall: 0.8431
F1 Score: 0.8687
ROC AUC: 0.9209
```

```
# Confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, xgb_preds), annot=True, fmt='d', cmap='Blues')
plt.title('XGBoost Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Feature importance
plt.figure(figsize=(10, 6))
xgb.plot_importance(xgb_model, max_num_features=15)
plt.title('XGBoost Feature Importance')
plt.show()
```

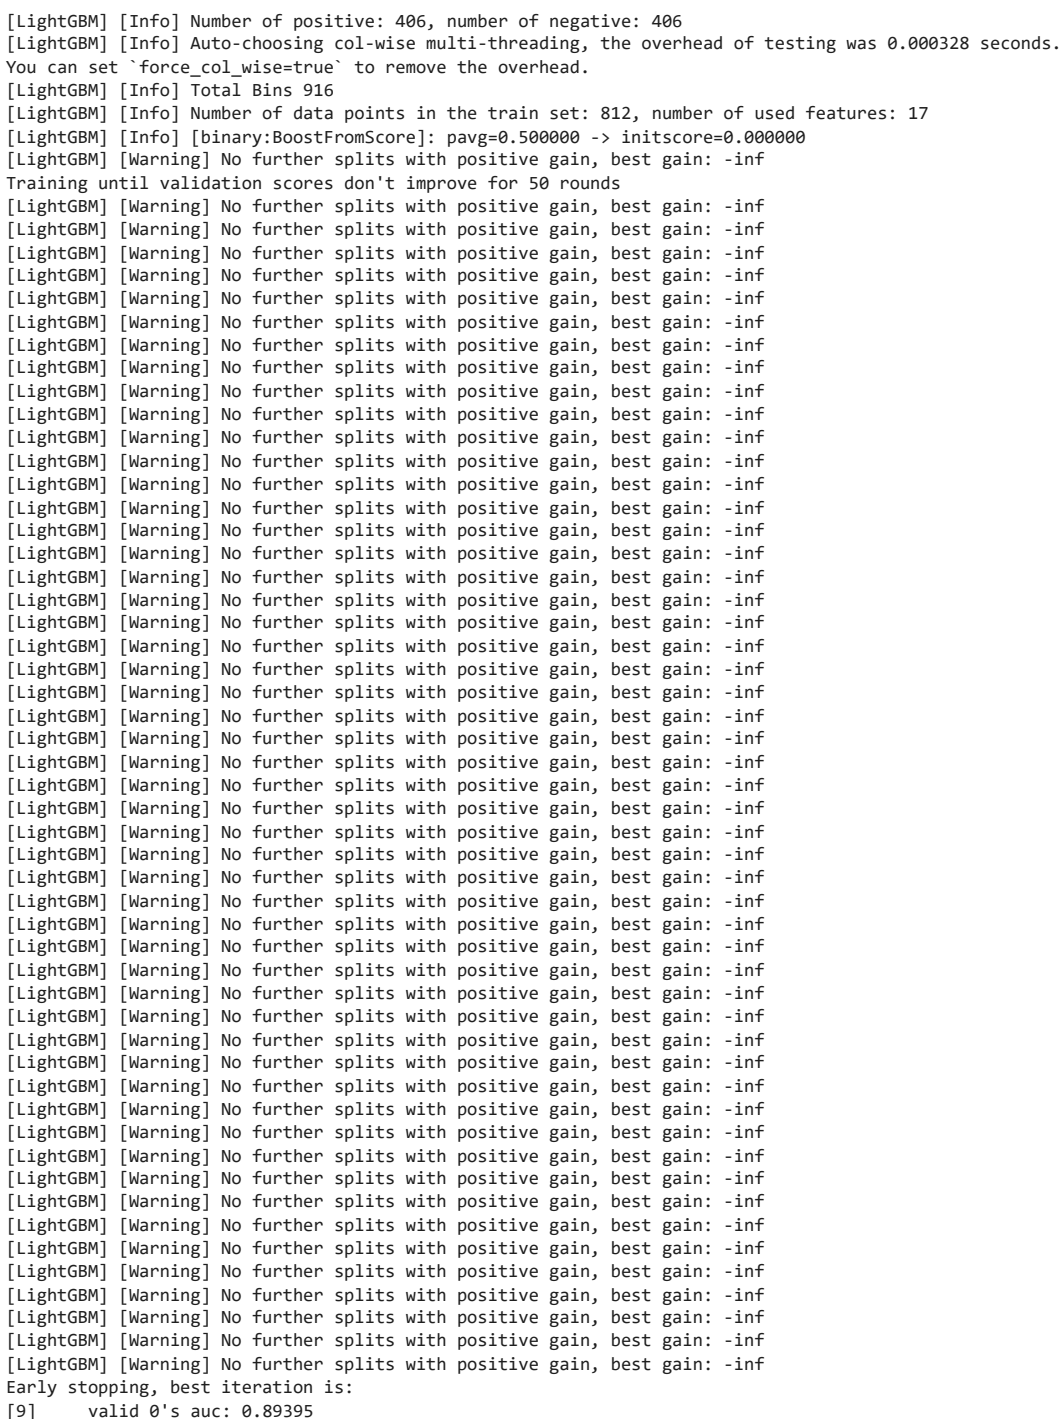
<Figure size 1000x600 with 0 Axes>



```
# LightGBM model
lgb_model = lgb.LGBMClassifier(
    objective='binary',
```


```
n_estimators=1000,  
learning_rate=0.01,  
max_depth=5,  
num_leaves=20,  
subsample=0.8,  
colsample_bytree=0.8,  
reg_alpha=0.1,  
reg_lambda=0.1,  
random_state=42,  
early_stopping_round=50,  
metric='auc'  
)
```

```
# Train the model  
lgb_model.fit(  
    X_train_preprocessed,  
    y_train_res,  
    eval_set=[(X_test_preprocessed, y_test)],  
    eval_metric='auc'  
)
```




```
LGBMClassifier(
    colsample_bytree=0.8, early_stopping_round=50,
    learning_rate=0.01, max_depth=5, metric='auc', n_estimators=1000,
    num_leaves=20, objective='binary', random_state=42,
    reg_alpha=0.1, reg_lambda=0.1, subsample=0.8)
```

```
# Evaluate performance
print("\nLightGBM Performance:")
print(f"Accuracy: {accuracy_score(y_test, lgb_preds):.4f}")
print(f"Precision: {precision_score(y_test, lgb_preds):.4f}")
print(f"Recall: {recall_score(y_test, lgb_preds):.4f}")
print(f"F1 Score: {f1_score(y_test, lgb_preds):.4f}")
print(f"ROC AUC: {roc_auc_score(y_test, lgb_probs):.4f}")
```




```
# Confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, lgb_preds), annot=True, fmt='d', cmap='Blues')
plt.title('LightGBM Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Feature importance
plt.figure(figsize=(10, 6))
lgb.plot_importance(lgb_model, max_num_features=15)
plt.title('LightGBM Feature Importance')
plt.show()
```

<Figure size 1000x600 with 0 Axes>



```
# Compare both models
results = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score', 'ROC AUC'],
    'XGBoost': [
        accuracy_score(y_test, xgb_preds),
        precision_score(y_test, xgb_preds),
        recall_score(y_test, xgb_preds),
        f1_score(y_test, xgb_preds),
        roc_auc_score(y_test, xgb_probs)
    ],
    'LightGBM': [
        accuracy_score(y_test, lgb_preds),
        precision_score(y_test, lgb_preds),
        recall_score(y_test, lgb_preds),
        f1_score(y_test, lgb_preds),
        roc_auc_score(y_test, lgb_probs)
    ]
})
```

```
print("\nModel Comparison:")
print(results)
```




```
Model Comparison:
Metric  XGBoost  LightGBM
0  Accuracy  0.858696  0.826087
1  Precision  0.895833  0.857143
2  Recall    0.843137  0.823529
3  F1 Score  0.868687  0.840000
4  ROC AUC   0.920911  0.893950
```

```
# Plot ROC curves
from sklearn.metrics import roc_curve
```

```
plt.figure(figsize=(8, 6))
fpr_xgb, tpr_xgb, _ = roc_curve(y_test, xgb_probs)
fpr_lgb, tpr_lgb, _ = roc_curve(y_test, lgb_probs)

plt.plot(fpr_xgb, tpr_xgb, label=f'XGBoost (AUC = {roc_auc_score(y_test, xgb_probs):.2f})')
plt.plot(fpr_lgb, tpr_lgb, label=f'LightGBM (AUC = {roc_auc_score(y_test, lgb_probs):.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison')
plt.legend()
plt.show()
```



```
# Create a pipeline with preprocessing and model
xgb_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', xgb.XGBClassifier(
        objective='binary:logistic',
        learning_rate=0.01,
        max_depth=5,
        random_state=42
    ))
])

lgb_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', lgb.LGBMClassifier(
        objective='binary',
        learning_rate=0.01,
        max_depth=5,
        random_state=42
    ))
])
```



```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=0.8, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               feature_weights=None, gamma=1, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=0.01, max_bin=None, max_cat_threshold=None,
               max_cat_to_onehot=None, max_delta_step=None, max_depth=5,
               max_leaves=None, min_child_weight=None, missing=nan,
               monotone_constraints=None, multi_strategy=None, n_estimators=1000,
               n_jobs=None, num_parallel_tree=None, ...)

```

```

# Save the model and preprocessing
import joblib
joblib.dump(final_model, 'final_heart_model.pkl')
joblib.dump(preprocessor, 'preprocessor.pkl')

```

```
['preprocessor.pkl']
```

```

# Example prediction function
def predict_heart_disease(input_data):
    """
    input_data: Dictionary with feature values
    Returns: Probability of heart disease
    """

    # Load model and preprocessor
    model = joblib.load('final_heart_model.pkl')
    preprocessor = joblib.load('preprocessor.pkl')

    # Convert input to DataFrame
    input_df = pd.DataFrame([input_data])

    # Apply feature engineering steps from the notebook
    # Create age bins (need to be consistent with training)
    input_df['AgeGroup'] = pd.cut(input_df['Age'], bins=[0, 40, 50, 60, 70, 100],
                                  labels=['<40', '40-50', '50-60', '60-70', '70+'])

    # Create HR/BP ratio
    input_df['HR_BP_Ratio'] = input_df['MaxHR'] / input_df['RestingBP']

    # Create cholesterol to age ratio
    input_df['Chol_Age_Ratio'] = input_df['Cholesterol'] / input_df['Age']

    # Convert Sex and ExerciseAngina
    input_df['Sex'] = input_df['Sex'].map({'M': 1, 'F': 0})
    input_df['ExerciseAngina'] = input_df['ExerciseAngina'].map({'Y': 1, 'N': 0})

    # One-hot encode other categorical features
    input_df = pd.get_dummies(input_df, columns=['ChestPainType', 'RestingECG', 'ST_Slope'], drop_first=True)

    training_columns = X_train.columns.tolist()

    # Add missing columns to input_df with default value 0 or False (depending on the type)
    for col in training_columns:
        if col not in input_df.columns:
            # Determine appropriate default value based on expected type after one-hot encoding
            # For one-hot encoded columns, default is False
            # For numeric/binary encoded columns, default is 0
            if col.startswith(('ChestPainType_', 'RestingECG_', 'ST_Slope_')):
                input_df[col] = False
            elif col == 'AgeGroup':
                pass
            else:
                input_df[col] = 0 # Default for numeric/binary

    # Ensure the columns are in the same order as the training data before dropping AgeGroup
    input_df = input_df[training_columns]

    # Drop AgeGroup from input_df as it was dropped from X_train_numeric before SMOTE and preprocessing
    input_df = input_df.drop('AgeGroup', axis=1)

```

```

# Now apply preprocessing
processed_data = preprocessor.transform(input_df)

# Make prediction
probability = model.predict_proba(processed_data)[0, 1]

return probability

# Example usage
sample_input = {
    'Age': 55,
    'Sex': 'M',
    'ChestPainType': 'ATA',
    'RestingBP': 130,
    'Cholesterol': 250,
    'FastingBS': 0,
    'RestingECG': 'Normal',
    'MaxHR': 150,
    'ExerciseAngina': 'N',
    'Oldpeak': 1.0,
    'ST_Slope': 'Flat'
}

print(f"\nPredicted probability of heart disease: {predict_heart_disease(sample_input):.2f}")

```



Predicted probability of heart disease: 0.72

```

low_risk_female = {
    'Age': 32,
    'Sex': 'F',
    'ChestPainType': 'NAP', # Non-Anginal Pain
    'RestingBP': 110,
    'Cholesterol': 180,
    'FastingBS': 0,
    'RestingECG': 'Normal',
    'MaxHR': 175,
    'ExerciseAngina': 'N',
    'Oldpeak': 0.0,
    'ST_Slope': 'Up'
}

print(f"\nPredicted probability of heart disease: {predict_heart_disease(low_risk_female):.2f}")

```



Predicted probability of heart disease: 0.44

```

high_risk_male = {
    'Age': 58,
    'Sex': 'M',
    'ChestPainType': 'TA', # Typical Angina
    'RestingBP': 145,
    'Cholesterol': 320,
    'FastingBS': 1,
    'RestingECG': 'ST',
    'MaxHR': 125,
    'ExerciseAngina': 'Y',
    'Oldpeak': 2.5,
    'ST_Slope': 'Flat'
}

print(f"\nPredicted probability of heart disease: {predict_heart_disease(high_risk_male):.2f}")

```



Predicted probability of heart disease: 0.97

```

moderate_risk = {
    'Age': 47,
    'Sex': 'M',
    'ChestPainType': 'ATA', # Atypical Angina
    'RestingBP': 130,
    'Cholesterol': 240,
    'FastingBS': 0,

```



```
'RestingECG': 'Normal',
'MaxHR': 150,
'ExerciseAngina': 'N',
'Oldpeak': 1.2,
'ST_Slope': 'Up'
}
print(f"\nPredicted probability of heart disease: {predict_heart_disease(moderate_risk):.2f}")
```



Predicted probability of heart disease: 0.58

```
elderly_female = {
    'Age': 72,
    'Sex': 'F',
    'ChestPainType': 'ATA',
    'RestingBP': 160,
    'Cholesterol': 210,
    'FastingBS': 0,
    'RestingECG': 'LVH', # Left Ventricular Hypertrophy
    'MaxHR': 115,
    'ExerciseAngina': 'N',
    'Oldpeak': 1.8,
    'ST_Slope': 'Flat'
}
print(f"\nPredicted probability of heart disease: {predict_heart_disease(elderly_female):.2f}")
```



Predicted probability of heart disease: 0.66

```
diabetic_case = {
    'Age': 55,
    'Sex': 'M',
    'ChestPainType': 'NAP',
    'RestingBP': 140,
    'Cholesterol': 190,
    'FastingBS': 1, # Diabetic
    'RestingECG': 'Normal',
    'MaxHR': 160,
    'ExerciseAngina': 'Y',
    'Oldpeak': 0.8,
    'ST_Slope': 'Up'
}
```